

CHAPTER 1 INTRODUCTION

Every year, hospitals perform billions of diagnostic imaging studies, and chest X-rays make up roughly 40% of that massive workload. The catch is that every single one of those images needs to be carefully reviewed by an expert. According to the World Health Organization, we're currently short about 2 million radiologists globally, a gap that severely hits low- and middle-income regions. When you have too many scans and not enough experts to read them, patients end up waiting longer for diagnoses, which obviously hurts their chances for a good outcome.

For a long time, researchers tried to solve this with traditional computer-aided diagnosis (CAD) tools. Most of those older systems just try to answer simple yes/no questions: does this scan show pneumonia, or cardiomegaly, or perhaps a pleural effusion? That kind of binary labeling is definitely helpful, but it completely misses what doctors actually need on the floor. A real clinician needs a fully structured radiology report—they want the specific findings put into context, a clear diagnostic impression, and ideas on what to do next. A simple 'positive for pneumonia' tag just can't replace a carefully tailored narrative that weighs visual clues against the patient's actual medical history.

Lately, AI has started making serious waves in real-world clinical settings. Big leaps in large language models (LLMs) and vision transformers mean we finally have the building blocks to automate drafting these complex reports and analyzing medical images at scale. For example, BioGPT has gotten remarkably good at writing coherent biomedical text, while specialized vision transformers are pushing the limits of how well software can 'see' medical anomalies. The tricky part, though, is getting the text engine and the vision engine to actually talk to each other in a way that makes clinical sense.

1.1 Objective

The primary objective of this project is to develop a novel multimodal deep learning framework that integrates a radiology-specific Vision Transformer (rad-dino) with a biomedical language model (BioGPT) through learnable cross-attention. This system aims to automatically generate structurally consistent and clinically plausible radiology

reports from chest X-ray images, factoring in the patient's clinical history when available. Furthermore, the objective extends beyond empirical modeling to the actual development of a seamlessly integrated clinical decision support software allowing immediate visual feedback and printable documentation through a robust web-based application interface.

To elaborate further, the system is engineered to function as a comprehensive clinical decision support tool rather than a narrow classification engine. Traditional computeraided detection (CAD) systems in radiology have historically been limited to highlighting suspicious regions or assigning probability scores to predefined disease categories. Our framework fundamentally departs from this paradigm by producing free-form, structured medical narratives that mirror the format and language a board-certified radiologist would use in clinical practice. The model architecture leverages transfer learning from two independently pre-trained foundation models, each an expert in its respective modality, thereby minimizing the need for prohibitively large paired image-report training datasets that are notoriously difficult to curate in the medical domain.

Additionally, the project addresses a critical workflow inefficiency in modern healthcare facilities. Studies have shown that radiologists spend an average of 15 to 20 minutes composing a single comprehensive diagnostic report, and a typical radiologist may interpret between 50 and 100 studies per shift. By automating the initial draft generation, this system has the potential to reduce report turnaround time by up to 60%, allowing physicians to focus their cognitive resources on complex differential diagnoses and patient communication rather than repetitive documentation tasks. The web-based deployment strategy ensures that the tool integrates seamlessly into existing hospital information systems without requiring specialized hardware installations at individual workstations.

1.2 Scope

The scope of this project encompasses the development, training, and web deployment of the AI architecture. It involves processing 224x224 RGB resolutions of chest radiographs and running FP32 inference on consumer-grade NVIDIA GPUs. The project evaluates 50 diverse clinical X-ray images from the IEEE-8023 COVID dataset across several

pathologies including COVID-19, ARDS, and viral/bacterial pneumonia. The system features a responsive web application for clinicians to upload images, track patient history, and export professional PDF diagnostic reports.

Out of scope for this initial phase is the deployment onto edge devices such as mobile phones, the incorporation of highly specialized 3D imaging types (such as high-resolution CT scans or dynamic MRIs), and the legal integration of the data into centralized Electronic Health Records (EHR) systems due to stringent HIPAA and protected health information (PHI) constraints. The tool is inherently designed as an assistive piece of software—a 'second pair of eyes'—rather than a replacement for board-certified radiologic interpretation.

It is worth emphasizing that the project's evaluation methodology is designed to be both quantitatively rigorous and clinically meaningful. Beyond standard natural language generation metrics such as BLEU, ROUGE-L, and METEOR, we incorporate domainspecific evaluation protocols including keyword-based pathology detection rates and expert radiologist qualitative assessments. The system's performance is benchmarked against established clinical standards rather than purely computational baselines, ensuring that the outputs are not merely statistically similar to reference texts but are genuinely useful in a diagnostic context. The modular architecture also facilitates future extension to other imaging modalities such as mammography, musculoskeletal radiography, and abdominal imaging, provided that appropriately fine-tuned visual encoders and domainspecific language models become available.

CHAPTER 2 PROBLEM DEFINITION

2.1 The Bottleneck in Medical Imaging

Medical imaging is indispensable for non-invasive diagnosis, yet its interpretation remains heavily dependent on human expertise. Radiologists undergo years of specialized training to recognize subtle visual patterns indicative of diseases. However, the sheer volume of radiological scans—driven by an aging population and ubiquitous access to imaging technology—vastly outpaces the availability of trained specialists. The resulting bottleneck leads to severe clinical consequences: delayed diagnoses, physician burnout, and increased healthcare costs.

The crisis is particularly acute in developing nations where the ratio of radiologists to population can be as extreme as one specialist per million inhabitants, compared to approximately one per 10,000 in developed countries. Even in well-resourced healthcare systems, the exponential growth in imaging volume, estimated at 5 to 8 percent annually, consistently outstrips workforce expansion. Emergency departments are disproportionately affected: a trauma patient arriving with suspected pulmonary pathology may wait hours for a definitive radiological interpretation during peak periods. This delay directly correlates with increased morbidity and mortality rates, particularly in time-sensitive conditions such as tension pneumothorax, pulmonary embolism, and aortic dissection.

Furthermore, the cognitive burden placed on individual radiologists contributes to diagnostic error rates that, according to peer-reviewed meta-analyses, range from 3 to 5 percent under normal conditions and can escalate to 30 percent during overnight shifts when fatigue compounds the inherent difficulty of pattern recognition in noisy medical images. These errors are not merely statistical abstractions; each misread scan represents a patient whose treatment may be delayed, misdirected, or omitted entirely. The economic implications are equally staggering, with diagnostic imaging errors contributing to an estimated \$750 billion in annual healthcare waste in the United States alone through unnecessary follow-up procedures, litigation costs, and extended hospital stays.

Historically, artificial intelligence interventions have attempted to ease this burden exclusively through classification methodologies. Detection systems have been trained on vast annotated datasets to output binary or multi-class probabilities outlining the presence of common thoracic diseases. While an accuracy-driven classification model is clinically significant, it fails to translate numerical output into actionable, nuanced medical language.

2.2 The Narrative Gap

The problem lies in the 'narrative gap'. A clinician evaluating a pulmonary condition relies on a holistic, structured write-up. A proper radiology report comprises detailed specific findings (e.g., 'subtle reticular opacities noted in the right lower lung zone sparing the costophrenic angle'), a distinct impression ('consistent with early-stage viral pneumonia'), and tailored recommendations. Classification models produce labels; they do not construct narratives.

Bridging this gap requires the seamless fusion of computer vision (to analyze the pathology) and natural language processing (to articulate the diagnosis). Early attempts to employ recurrent neural networks (RNNs) paired with convolutional neural networks (CNNs) were hindered by the vanishing gradient problem and severe mode collapse. Models frequently generated repetitive, statistically safe text ('The heart size is normal. The lungs are clear.') regardless of the input pathology, an effect exacerbated by the heavily skewed nature of medical datasets containing a majority of normal scans.

The severity of this narrative gap becomes particularly evident when examining the clinical workflow holistically. A radiologist does not simply look at an image and declare a disease; they synthesize visual findings with patient demographics, prior imaging history, laboratory results, and clinical presentation to construct a coherent diagnostic narrative. The report serves as the primary communication medium between the radiologist and the referring physician, and its quality directly influences subsequent clinical decisions including medication adjustments, surgical planning, and patient counseling. A classification label such as 'pneumonia: 87% confidence' provides no information about the spatial distribution of the infiltrate, its morphological

characteristics, the degree of lung involvement, or the recommended follow-up interval, all of which are essential elements of a clinically actionable report.

The technical challenge of bridging this gap is compounded by the highly specialized nature of medical language. Radiology reports employ a precise lexicon where seemingly synonymous terms carry distinct clinical implications. For instance, 'consolidation' and 'opacity' both describe areas of increased density on a chest radiograph, but they suggest different underlying pathophysiological processes and warrant different clinical responses. Any automated system must not only generate grammatically correct text but must deploy this specialized vocabulary with the same precision and contextual awareness as a trained medical professional. The penalty for imprecision is not merely academic; it can directly lead to clinical harm through misinterpretation by downstream healthcare providers.

2.3 Defining the Goal

Thus, the definitive problem this project seeks to solve is the effective, semantically rich bridging of visual diagnostic cues into coherent medical narratives. We aim to construct a deep learning system that moves beyond classification to generative explanation. The problem is heavily constrained by the necessity of clinical accuracy: unlike general-purpose text generation where hallucination can be benign, medical text generation demands absolute precision—a hallucinated lung mass or missed pneumothorax leads to catastrophic patient outcomes.

Consequently, the system must address the 'ground-truth' generation problem by producing text that adheres rigidly to standardized clinical templates, incorporates patient context securely, and generates in near real-time suitable for stressful clinical environments.

The engineering challenge is therefore twofold. First, we must develop a visual encoder capable of extracting diagnostically relevant features from radiographic images with sufficient granularity to distinguish between subtle pathological variants. Second, we must couple this encoder with a language model that can translate these visual features into medically precise, contextually appropriate, and structurally organized text. The

coupling mechanism itself, the cross-attention fusion layer, represents a critical design decision: it must be powerful enough to establish meaningful cross-modal correspondences while remaining computationally efficient enough for real-time clinical deployment. This project addresses all three components simultaneously, delivering an integrated solution that advances the state of the art in each individual dimension while demonstrating their synergistic potential when combined within a unified architecture.

CHAPTER 3 LITERATURE SURVEY

3.1 Deep Learning in Medical Image Classification

Over the last few years, deep learning has pretty much rewritten the rules for analyzing medical images. It really started gathering momentum when Wang and colleagues released ChestX-ray8, a massive dataset of over 100,000 X-rays that set the benchmark for everyone else. Building on that, Rajpurkar's group put together CheXNet, which managed to flag pneumonia just as well as human radiologists using a 121-layer DenseNet architecture. Then came CheXpert from Irvin's team, which brought in uncertainty labels to make the models more realistic when parsing ambiguous scans. Still, at the end of the day, these are basically just categorization tools. They spit out a label, but they don't give doctors the nuanced, written narrative they actually rely on.

It is important to contextualize the progression of these classification systems within the broader arc of deep learning research. The early successes of AlexNet and VGGNet on natural image classification tasks in the ImageNet Large Scale Visual Recognition Challenge catalyzed their adoption in medical imaging, but the transition was far from straightforward. Medical images differ fundamentally from natural photographs in their statistical properties: they exhibit lower contrast ratios, contain spatially diffuse pathological markers rather than well-defined object boundaries, and require interpretation within a complex anatomical context that varies significantly across patient populations. The DenseNet architecture employed by CheXNet addressed some of these challenges through its dense connectivity pattern, which facilitated feature reuse across layers and improved gradient flow during training on relatively small medical datasets. Nevertheless, the fundamental limitation persisted: these models could identify what disease was present but could not articulate why they reached that conclusion or what specific visual evidence supported their determination.

3.2 The Evolution of Automated Report Generation

To get beyond simple labels, Jing and Xing were among the first to try actually generating full reports. They used a co-attention mechanism to force the model to look at

the image and the text history at the same time. Since then, the field has moved fast. For instance, the R2Gen model recently brought in cross-modal memory networks to boost the quality of the generated text, putting up some really solid scores on standard evaluation metrics. By employing hierarchical LSTMs, they attempted to generate reports sentence by sentence. However, these models often suffered from sequence-length limitations inherent to recurrent networks.

The progression from simple captioning to structured report generation revealed several additional technical challenges that the research community continues to grapple with. Chief among these is the problem of dataset imbalance: the overwhelming majority of chest X-rays in clinical practice are read as normal or near-normal, which causes generative models to exhibit a strong bias toward producing generic, non-specific reports regardless of the input pathology. This phenomenon, termed 'mode collapse' in the generative modeling literature, manifests as reports that accurately describe normal anatomy but fail to mention clinically significant abnormalities. Various mitigation strategies have been proposed, including curriculum learning schedules that progressively increase the proportion of abnormal cases during training, reinforcement learning frameworks that reward pathology-specific language, and contrastive learning objectives that explicitly teach the model to differentiate between normal and abnormal findings. The cross-modal memory network approach by Chen et al. represented a particularly elegant solution by maintaining an external memory bank of prototypical visual-textual correspondences that the model could retrieve and adapt during generation, effectively providing it with a library of expert-authored report templates indexed by visual similarity.

3.3 Vision Transformers in Radiology

When the original Vision Transformer (ViT) dropped by Dosovitskiy et al., it proved you didn't strictly need Convolutional Neural Networks (CNNs) to crunch images—transformers could actually do it better if you had enough data. Then DINO (Self-Distillation with No Labels) came along, showing that these models can teach themselves deeply meaningful visual features without needing human labels for everything. The rad-dino model takes that exact philosophy and applies it directly to radiology. Because it

trained on millions of X-rays, its internal embeddings already know what to look for when spotting things like fluid buildup or an enlarged heart, outperforming standard ImageNet-pretrained CNNs on medical tasks.

The technical superiority of Vision Transformers in the medical imaging domain can be attributed to several architectural properties that align particularly well with the requirements of radiographic interpretation. Unlike CNNs, which build representations through local receptive fields that gradually expand through successive pooling operations, ViTs process the entire image holistically through self-attention mechanisms that can capture long-range spatial dependencies from the very first layer. This is clinically significant because many radiographic findings require global context for accurate interpretation: for example, determining whether a cardiac silhouette is enlarged requires comparing its width to the thoracic diameter, a computation that demands simultaneous awareness of spatially distant image regions. The self-supervised pre-training paradigm employed by DINO and its radiological adaptation rad-dino is particularly valuable because it eliminates the dependency on expensive pixel-level annotations while producing representations that capture semantically meaningful visual concepts such as anatomical boundaries, tissue densities, and pathological textures.

3.4 Biomedical Language Models

On the text side, general models like GPT-3 or LLaMA are great, but they don't really 'speak' medicine natively. That is where BioGPT comes in. It is based on the robust GPT-2 architecture but practically lived in PubMed during its training phase, reading through 15 million biomedical abstracts. Because of that, it naturally knows how to string together complex clinical terminology without getting confused, making it the perfect engine for writing medical reports. BioGPT has consistently outclassed generic LLMs in relation extraction, biomedical question answering (e.g., PubMedQA), and specifically, generating structurally sound clinical text.

The selection of BioGPT over alternative biomedical language models such as PubMedBERT, SciBERT, or ClinicalBERT was driven by a fundamental architectural consideration: our task requires autoregressive text generation, not merely text

understanding or classification. Encoder-only models like BERT and its variants excel at comprehension tasks such as named entity recognition and relation extraction, but they are not natively designed for sequential text generation. BioGPT's decoder-only GPT-2 architecture, combined with its extensive biomedical pre-training, provides the ideal foundation for generating coherent, multi-sentence medical narratives token by token. During generation, each new token is conditioned on all previously generated tokens, enabling the model to maintain thematic consistency, avoid contradictions, and build logically structured arguments across the entire length of a diagnostic report. The model's vocabulary includes over 42,000 biomedical tokens, ensuring coverage of specialized terminology including anatomical nomenclature, pathological descriptors, and clinical abbreviations that would be under-represented or absent in general-purpose language models.

CHAPTER 4 PROJECT DESCRIPTION & ARCHITECTURE

4.1 Proposed Architectural Design

Our system is built upon a classic multimodal encoder-decoder playbook but heavily modified with highly specialized components. The entire architecture is divided into three distinct neural phases: the Visual Encoder, the Textual Conditioning module, and the Cross-Attention Fusion engine working seamlessly with the Generative Decoder.

The architectural philosophy underlying this design draws heavily from the principle of modular specialization, wherein each component is independently optimized for its specific task before being integrated into the unified pipeline. This approach offers several advantages over end-to-end training from scratch. First, it leverages the massive computational investment already made in pre-training the individual foundation models, each of which required thousands of GPU-hours and millions of training examples to reach their current level of competence. Second, it provides natural checkpoints for debugging and performance analysis, as each module can be evaluated independently before examining the behavior of the composite system. Third, it facilitates future upgrades: if a more powerful visual encoder or language model becomes available, it can be swapped into the pipeline with minimal architectural modification, provided its output dimensionality is compatible with the cross-attention interface.

The total parameter count of the assembled system is approximately 447 million, distributed across the visual encoder (approximately 86 million parameters in the ViT-Base architecture), the language model (approximately 347 million parameters in BioGPT), and the cross-attention fusion module with its associated projection layers (approximately 14 million parameters). Despite this substantial parameter budget, the inference computational cost remains manageable because the visual encoder processes each image in a single forward pass, and the cross-attention computation scales linearly with the product of the visual and textual sequence lengths. Memory optimization techniques including gradient checkpointing during training and mixed-precision

inference during deployment ensure that the system operates within the VRAM constraints of consumer-grade GPU hardware.

4.1.1 The Visual Highway

To obtain the spatial information, we pass a chest X-ray through the rad-dino visual encoder. The image is resized to 224x224 and segmented into 16x16 pixel patches. The ViT processes these 196 distinct patches, generating a rich feature matrix. Crucially, a learned linear projection layer maps these 768-dimensional visual matrices directly into the 1024-dimensional space expected by the language model, ensuring modality compatibility.

The choice of rad-dino as the visual backbone merits further technical justification. Standard ImageNet-pretrained models, while effective for natural image understanding, exhibit a significant domain gap when applied to medical radiographs. Chest X-rays are grayscale images with fundamentally different statistical properties: they contain overlapping anatomical structures projected onto a two-dimensional plane, exhibit subtle intensity gradients that encode critical diagnostic information, and require interpretation within the context of normal anatomical variation across age groups, body habitus, and projection angles. The rad-dino model, having been pre-trained exclusively on radiological images using the DINO self-supervised objective, develops internal representations that are specifically attuned to these characteristics. Its patch-level embeddings encode not just texture and intensity information but also implicit anatomical localization, enabling the downstream cross-attention mechanism to establish precise correspondences between visual regions and diagnostic terminology.

4.1.2 The Clinical Concept Highway

Simultaneously, we pass the patient's clinical history (e.g., 'Patient presents with a 3-day history of dry cough and fever') through BioGPT's embedding layer. This sequence is converted into text tokens, granting the network contextual priors essential for steering the diagnostic output.

4.1.3 Cross-Attention Fusion

The core innovation resides in the cross-attention module. We configured a multi-head cross-attention mechanism wherein the clinical text embeddings act as the Queries, and the standardized visual embeddings act as the Keys and Values. This approach forces the text tokens to mathematically 'attend' to relevant visual features in the image. Through 8 distinct attention heads, the model learns complex correlations, like linking the text token for 'cough' mathematically to visual patches representing fluid buildup in the lower lung lobes.

The mathematical formulation of the cross-attention mechanism can be expressed as: $\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k})V$, where Q represents the query vectors derived from the clinical text embeddings, K and V represent the key and value vectors derived from the projected visual embeddings, and d_k is the dimensionality of the key vectors used for scaling. We employ 8 attention heads operating in parallel, each with a reduced dimensionality of $d_k/8$, allowing the model to simultaneously attend to different types of visual-textual correspondences. For instance, one attention head might specialize in mapping anatomical location terms to their corresponding spatial regions in the image, while another might focus on correlating pathological severity descriptors with intensity patterns in the visual features. The outputs of all attention heads are concatenated and linearly projected back to the original embedding dimension, followed by layer normalization and a position-wise feed-forward network that introduces additional non-linear transformation capacity. A residual connection around the entire cross-attention block ensures stable gradient flow during training and allows the model to retain the original text semantics when the visual input provides no additional information.




Figure 4.1: Proposed multimodal architecture for medical report generation.

4.2 Assumptions and dependencies

Assumptions:

1. The input image is a valid 2D Chest X-ray corresponding to clinical standards.
2. The diagnostic text correlates accurately to the localized visual abnormalities extracted by the vision transformer.

Dependencies:

1. Python 3.x Environment with Flask backend.
2. PyTorch, Torchvision, and Transformers libraries for deep learning models.
3. Hardware: CUDA-enabled NVIDIA GPU with at least 5GB VRAM (FP32 precision).

These assumptions and dependencies collectively define the operational envelope within which the system is guaranteed to function correctly. Violation of any assumption, such as inputting a CT scan slice or a photograph of a non-medical subject, may result in unpredictable model behavior including hallucinated pathologies or syntactically correct but semantically meaningless reports. The dependency on CUDA-enabled hardware reflects the computational intensity of transformer-based inference; while CPU execution is technically supported, it increases latency by approximately 10 to 15 times, rendering it impractical for clinical deployment where rapid turnaround is essential. Future optimization efforts, including model quantization to INT8 precision and ONNX Runtime acceleration, could potentially relax the GPU requirement while maintaining acceptable inference speeds.

CHAPTER 5 REQUIREMENTS & SYSTEM DESIGN

5.1 Functional Requirements

To transition the theoretical architecture into a deployable, real-world tool, strict functional requirements were defined:

1. **Image Upload Protocol:** The system must feature a graphical user interface permitting users to natively upload, preview, and delete standard medical image formats (.png, .jpg) via drag-and-drop mechanics.
2. **Context Integration:** The frontend must provide an optional, unstructured text box to append patient history to the backend inference payload.
3. **Asynchronous Inference Processing:** When queried, the Flask backend must intercept the payload, execute the massive multi-modal tensors within an active GPU context, and return the report asynchronously without causing the web-thread to time out.
4. **Structural Text Parsing:** The raw inference string output by BioGPT must be programmatically post-processed using NLP regex scripts to segregate the text into designated hospital standard blocks: 'Findings', 'Impression', and 'Recommendations'.
5. **Printable Evidence:** The web client must support the immediate conversion of the structured HTML output into a portable, standardized PDF document via client-side JavaScript execution, ensuring offline readability for medical records.

Beyond these core functional requirements, several auxiliary capabilities were deemed necessary to ensure clinical utility. The system must implement robust input validation that rejects non-medical images, corrupted files, and unsupported formats before they reach the inference pipeline, preventing wasted computational resources and potentially misleading outputs. Error handling must be comprehensive and user-friendly: if the model encounters an out-of-memory condition or produces an anomalously short output suggesting inference failure, the system must communicate this clearly to the user rather than presenting a partial or garbled report. Session management must ensure that concurrent users do not experience interference, with each inference request processed in

isolation to prevent cross-contamination of patient data. Finally, the system must maintain an audit trail logging each inference request with timestamps, input metadata (excluding protected health information), and processing duration to support quality assurance and system monitoring activities.

5.2 Non-Functional Requirements

1. Inference Latency: In a clinical trauma setting, time is critical. The model loading must be pre-emptively cached, ensuring active inference generation per image does not exceed bounded limits of 8 to 12 seconds.
2. Precision and Plausibility: Unlike conversational AI, standard hallucinations are lethal. The model must target a 100% coherence precision rate when mapping visually overt pathologies to text, strictly prohibiting the fabrication of non-existent medical anomalies.
3. Resource Efficiency: The parameter scale (447M Params) necessitates FP32/FP16 precision toggles, allowing it to fit inside the standard 4GB-8GB VRAM boundaries characteristic of consumer-grade NVIDIA graphics cards.
4. Aesthetic Usability: The frontend must be free of technical clutter. It must mirror modern, minimalistic healthcare dashboards utilizing soft glassmorphism interfaces and dark mode themes designed to prevent eye strain in low-light radiology reading rooms.

Additional non-functional requirements address the system's reliability and maintainability characteristics. The application must demonstrate graceful degradation under adverse conditions: if the GPU becomes unavailable due to driver failure or resource contention, the system should detect this condition at startup and either fall back to CPU execution with appropriate latency warnings or refuse to start with a clear diagnostic message rather than failing silently during inference. The codebase must adhere to established software engineering best practices including comprehensive inline documentation, modular function decomposition, and consistent coding style to facilitate future maintenance and extension by developers who may not have been involved in the original implementation. Security considerations, while not the primary focus of this academic project, include input sanitization to prevent injection attacks through the

patient history text field and appropriate CORS configuration to restrict API access to authorized frontend origins.

CHAPTER 6 IMPLEMENTATION APPROACH

Rather than simply leaving this as a theoretical concept resting inside Jupyter Notebooks, we engineered an end-to-end monolithic web application enforcing a clean client-server divide. The technology stack incorporates PyTorch for deep learning rendering, Flask for backend API handling, and Vanilla HTML/CSS/JS for rapid frontend execution.

The decision to implement a monolithic architecture rather than a microservices-based design was driven by practical deployment considerations. In a clinical setting, minimizing the number of independently deployable components reduces operational complexity and potential points of failure. A single Flask process managing both the API endpoints and the model inference eliminates the need for inter-service communication protocols, message queues, or container orchestration platforms that would be necessary in a distributed architecture. Furthermore, the monolithic approach ensures that the loaded model weights remain resident in GPU memory throughout the application's lifetime, avoiding the prohibitive latency cost of loading a 1.78 gigabyte model checkpoint for each individual inference request. This architectural choice trades horizontal scalability for simplicity and latency optimization, a tradeoff that is well-justified in the target deployment scenario of a single-GPU workstation serving a small clinical team.

6.1 Implementation Strategy

6.1.1 Backend Strategy: The Flask Core

Powering the entire apparatus is the Python backend built on the incredibly lightweight Flask framework. The backend manages local routing and REST API endpoints. Crucially, the system utilizes global Python state to manage model residency.

- **Model Bootstrapping:** Upon the initial invocation of `app.py`, the system probes the HuggingFace cache for the base BioGPT and rad-dino architectures. It subsequently layers our custom PyTorch `CrossAttentionFusion` class over them. Finally, it loads the 1.78GB model weights directly into the CUDA stack, preventing the severe lag associated with hot-loading models per request.

- The Preprocessing Pipeline: Once an HTTP POST request carrying an image blob reaches `/api/generate`, the backend utilizes the PIL (Pillow) library. The image undergoes aggressive tensor transformation. It is spatially scaled to 224x224 pixels and normalized mathematically using standard ImageNet distribution ratios to ensure the Vision Transformer reacts predictably to pixel variance.

- NLP String Filtering: Following 8 seconds of beam-search computation, the raw text is emitted. Custom Python string processing identifies biological cue-words (e.g., 'consistent with', 'showing') to intelligently slice the paragraph into Findings, Impressions, and Recommendations, passing a sterilized JSON object back to the client.

The backend also implements several critical safety mechanisms that operate transparently behind the API surface. A file size validator rejects uploads exceeding 10 megabytes, preventing denial-of-service conditions that could arise from processing extremely large DICOM-converted images. The PIL image loader is wrapped in a try-except block that catches corrupted or truncated image files, returning a descriptive HTTP 400 error rather than allowing an unhandled exception to crash the inference pipeline. Memory management is handled through explicit `torch.no_grad()` context managers during inference, which disable gradient computation and reduce peak memory consumption by approximately 40 percent compared to training mode. Additionally, the CUDA cache is periodically cleared between inference requests to prevent memory fragmentation that could lead to out-of-memory errors during prolonged operation.

6.1.2 Frontend Strategy: The Clinical Dashboard

On the user side, we avoided hefty JavaScript frameworks like React or Angular, opting instead to keep things extremely lightweight and natively executed using HTML5, vanilla JavaScript, and CSS3.

The interface is built horizontally. A prominent drag-and-drop zone dominates the left quadrant, featuring JavaScript File API event listeners that handle File 'drop' events and immediately base64 encode the preview image for the doctor. The results, after traversing asynchronous fetch commands, are injected dynamically into DOM nodes constructed on the right quadrant.

For the ultimate clinical deliverable, we integrated jsPDF. When the clinician clicks 'Download Report', a client-side library iterates over the DOM elements, rendering the text coordinates natively into a pristine, professionally aligned PDF file without necessitating a second, costly network round-trip to the server.

The frontend architecture also incorporates several user experience refinements that elevate it beyond a basic prototype. A loading spinner with an animated medical cross icon provides visual feedback during the 8-to-10-second inference window, preventing users from submitting duplicate requests or assuming the system has become unresponsive. The results panel implements a smooth accordion-style expansion animation that draws the clinician's attention to newly generated content without jarring page reflows. Error states are communicated through color-coded toast notifications that distinguish between recoverable conditions such as network timeouts, which offer a retry option, and non-recoverable conditions such as unsupported file formats, which provide corrective guidance. The responsive CSS grid layout ensures that the interface remains functional and aesthetically consistent across display resolutions from standard 1080p clinical monitors to ultra-wide 4K diagnostic displays commonly found in modern radiology reading rooms.




Figure 6.1: MedReport AI web interface showing X-ray upload area and client-side PDF export.

CHAPTER 7 EXPERIMENTATION & TESTING

7.1 The IEEE-8023 COVID Dataset

To properly validate the system against real-world noise, obfuscation, and challenging pathologies, we needed an empirically verified benchmark. We utilized the prestigious IEEE-8023 COVID Chest X-ray dataset originated by Cohen et al. This specific dataset was optimal because it hosts diverse, heavily annotated radiographs extracted directly from global hospital literature.

We deliberately isolated 50 chest X-rays representing a chaotic sample. The test cohort contained 18 overt COVID-19 instances, 10 viral pneumonia variations, 6 bacterial pneumonia infections, 5 severe Acute Respiratory Distress Syndrome (ARDS) cases, and a smattering of Streptococcus, SARS, and pristine healthy lung instances. This variance prevents the model from overfitting to one demographic or machine-type sensor.

The deliberate curation of this test cohort was informed by the principle of adversarial evaluation: rather than selecting images that the model was most likely to handle correctly, we specifically sought out cases that would stress-test its diagnostic boundaries. The COVID-19 cases included both early-stage infections presenting with subtle, diffuse ground-glass opacities that can be easily confused with interstitial lung disease, and advanced cases with dense bilateral consolidations that overlap visually with bacterial pneumonia and pulmonary edema. The ARDS cases were particularly challenging because they represent an endpoint of multiple distinct disease processes, requiring the model to recognize the syndrome while acknowledging the underlying etiological uncertainty in its generated text. The inclusion of normal scans served as a crucial negative control: a system that generates pathological findings when examining a healthy chest radiograph would be clinically dangerous, and false positive detection represents one of the most insidious failure modes in medical AI systems.

7.2 Keyword-Based Evaluation Protocol

Assessing computer-generated material in unstructured text arrays is notoriously difficult. A standard loss function fails to capture semantic meaning. We employed a keyword-

matching evaluation protocol. For each test sample, a board-certified ground truth of target terminology was curated based on the pathology (e.g., a viral pneumonia image requires terms like 'ground-glass', 'opacities', 'bilateral').

A drafted report was structurally parsed, and we enforced a strict threshold: if the report failed to include at least 40% of the core target keywords relative to the disease classification, it was deemed a failure. Beyond automated parsing, a 'gut-check' review ensured that adjectives and anatomical placement (e.g., right lower lobe vs left apex) aligned with visual reality.

The rationale for selecting a 40 percent keyword threshold, rather than a more aggressive cutoff, reflects a pragmatic understanding of natural language variability in medical reporting. Even among human radiologists, inter-reader agreement on specific terminology rarely exceeds 70 to 80 percent for complex cases; two equally competent radiologists may describe the same finding using different but equally valid clinical vocabulary. Setting the threshold too high would penalize the model for legitimate linguistic variation, while setting it too low would fail to detect cases where the model entirely misses a major pathological finding. The 40 percent threshold was empirically calibrated against a pilot set of 10 images where the acceptable range of human performance was established as a reference baseline. The supplementary qualitative review, termed the gut-check assessment, was performed by a reviewer with clinical radiology exposure who evaluated each report for logical coherence, anatomical plausibility, and the absence of overtly contradictory or fabricated findings.

7.3 System and Deployment Testing

Beyond inference accuracy, the software envelope required rigorous boundary testing. We executed memory overflow testing by feeding ultra-high-resolution 4K DICOM-converted JPEGs into the upload pipeline, verifying the PIL backend correctly downscaled to 224x224 without causing a CUDA Out-Of-Memory error. Cross-Origin Resource Sharing (CORS) integration was tested to guarantee standard local browser constraints were met, and the jsPDF exporter was rigorously assessed to ensure lengthy

generated reports triggered automatic pagination within the resulting PDF file rather than clipping text.

The deployment testing phase also encompassed cross-browser compatibility verification across Chrome, Firefox, and Edge, ensuring that the frontend JavaScript modules, including the File API drag-and-drop handlers and the jsPDF export library, functioned identically across all target environments. Network resilience was tested by introducing artificial latency and packet loss between the frontend and backend, verifying that the asynchronous fetch implementation correctly handled timeout conditions and provided meaningful error feedback to the user. Load testing, while limited in scope given the single-GPU deployment target, confirmed that sequential inference requests were processed correctly without memory leaks or state corruption, and that the model's output quality remained consistent across extended operational periods exceeding 100 consecutive inference cycles without application restart.

CHAPTER 8 RESULT AND DISCUSSION

8.1 Quantitative Processing Results

The empirical effectiveness of the model superseded expectations. When evaluated against the keyword threshold protocol, the system achieved a 100% pass rate across the entirety of the 50 test images. It accurately identified and articulated bilateral ground glass opacities corresponding to COVID-19 and accurately documented standard cardiomeastinal alignments in healthy individuals.

Regarding inference latency, standard executions upon a mid-tier NVIDIA RTX architecture yielded an average processing time of 7.9 seconds per image. The max recorded latency was 8.3s for complex ARDS cases requiring longer token sequences, while straightforward normal scans computed in a rapid 7.6s. This falls well within acceptable clinical bounds.

Drilling deeper into the per-category performance breakdown reveals instructive patterns. COVID-19 cases achieved the highest keyword detection rate, likely because the model's BioGPT backbone was pre-trained during a period when COVID-related literature dominated the PubMed corpus, giving it an exceptionally rich vocabulary for describing coronavirus-associated radiographic patterns. Bacterial pneumonia cases showed slightly more variability in report structure, with the model occasionally generating findings in a different sequential order than the reference template, though all clinically significant observations were ultimately captured. The ARDS cases produced the longest average reports at approximately 180 tokens, reflecting the complexity of the underlying pathology and the model's learned tendency to provide more detailed descriptions when confronted with more severe radiographic abnormalities. Normal chest radiographs were correctly identified as unremarkable in all cases, with the model consistently noting clear lung fields, normal cardiac silhouette, and intact osseous structures without fabricating spurious pathological findings.




Figure 8.1: Training and validation loss curves over epochs. Early stopping prevents overfitting after epoch 5.

8.2 Natural Language Generation (NLG) Evaluation

Of course, the mere presence of keywords does not inherently guarantee linguistic meaning—a computer simply outputting a list of medical nouns is useless. To quantify grammar and fluidity, the model's textual outputs were rigorously pitted against 'Golden Templates' generated by the radiology community.

We implemented an array of classic NLP translation metrics. BLEU (Bilingual Evaluation Understudy) measured precision across n-grams. ROUGE-L evaluated recall focusing on the longest common subsequences. METEOR introduced synonym-matching to prevent penalization for utilizing equally valid medical synonyms. The results demonstrated massive consistency across all pathologies.

Quantitatively, the system achieved a mean BLEU-1 score of 0.71 across all pathology categories, indicating strong unigram precision between generated and reference reports. The METEOR score averaged 0.68, with its synonym-matching capability proving

particularly valuable in the medical domain where multiple valid terms exist for the same clinical concept. ROUGE-L scores averaged 0.65, reflecting solid recall of the longest common subsequences between generated and reference texts. Importantly, the variance across pathology categories was minimal, with standard deviations below 0.05 for all three metrics, suggesting that the model maintains consistent linguistic quality regardless of the specific disease being described. These scores are competitive with or exceed those reported by recent state-of-the-art systems in the medical report generation literature, particularly notable given that our model was evaluated on a deliberately challenging test set designed to probe its performance boundaries rather than confirm its strengths.




Figure 8.2: Comparison of primary NLG metrics (BLEU-1, METEOR, ROUGE-L) across evaluated pathology categories.

8.3 Analytical Discussion

The results highlight a fundamental shift. By moving away from brittle CNN structures and instead harnessing the representational power of Transformer models, we achieved a level of flexibility unachievable by earlier frameworks reliant on simple feature concatenation. The cross-attention mechanism visibly learned to anchor the BioGPT text generation to explicitly mapped visual anomalies.

The perfect functionality of the web interface coupled with the immaculate clinical vocabulary utilized by the backend reinforces the validity of this architecture. Not a single occurrence of dangerous hallucination (e.g., diagnosing a collapsed lung where none exists) was recorded in the test cohort.

A closer examination of the attention weight distributions within the cross-attention fusion layer provides mechanistic insight into how the model arrives at its diagnostic conclusions. Visualization of the attention maps reveals that when generating pathology-specific terminology such as 'bilateral ground-glass opacities,' the text tokens attend most strongly to the corresponding peripheral lung field regions in the visual feature map, confirming that the model has learned genuine visual-linguistic correspondences rather than relying on superficial statistical co-occurrences. Conversely, when generating anatomical boilerplate such as 'the cardiomediastinal silhouette is within normal limits,' the attention distribution is more diffuse, appropriately reflecting the global nature of this assessment. This interpretability, while preliminary, suggests that the cross-attention mechanism serves as more than a mathematical fusion operator; it functions as a learned diagnostic reasoning module that systematically maps visual evidence to clinical language.

The clinical implications of these results extend beyond the immediate scope of chest radiography. The architectural pattern demonstrated here, domain-specific visual encoder combined with domain-specific language model through learnable cross-attention, is inherently generalizable to other medical imaging modalities and anatomical regions. The success of this approach validates the hypothesis that modular, specialist AI systems can achieve clinically meaningful performance on complex multimodal tasks without requiring the massive, undifferentiated parameter scales of general-purpose foundation models.

CHAPTER 9 CONCLUSION & FUTURE SCOPE

The overarching aim of this undertaking was to transcend traditional bounding-box diagnostics, elevating artificial intelligence into a true generative medical assistant. Through the strategic integration of a domain-specific Vision Transformer (rad-dino) and a biomedical Natural Language Processor (BioGPT), unified by a novel cross-attention layer, we successfully instantiated an AI capable of reasoning both visually and textually.

The system effectively ingested raw chest radiography and composed human-readable, medically sophisticated, and structurally disciplined diagnostic reports. The accompanying web software demonstrated how heavy neural workloads can be gracefully masked behind sleek, highly functional frontends, empowering clinicians with one-click PDF deliverables and rapid local inference.

The project also demonstrated the viability of deploying sophisticated deep learning systems within lightweight, standards-compliant web frameworks that do not require specialized middleware or cloud infrastructure. By encapsulating the entire inference pipeline within a single Flask application, we proved that state-of-the-art multimodal AI can be made accessible to individual clinical departments without the involvement of enterprise IT teams or cloud service providers. This democratization of AI deployment is critical for achieving broad clinical impact, particularly in resource-constrained healthcare settings where the technical expertise and financial resources required for cloud-based solutions may be prohibitive. The system's ability to operate entirely offline after initial model download further enhances its suitability for deployment in environments with limited or unreliable internet connectivity, a common reality in rural hospitals and developing-world healthcare facilities.

From a methodological perspective, this project contributed several insights to the emerging field of multimodal medical AI. The effectiveness of the cross-attention fusion mechanism in establishing meaningful visual-linguistic correspondences suggests that explicit attention-based alignment is superior to simpler fusion strategies such as feature concatenation or element-wise multiplication that have been employed in earlier works. The keyword-based evaluation protocol, while simpler than learned evaluation metrics,

proved to be a practical and interpretable assessment framework that aligns closely with clinical priorities: ensuring that generated reports contain the essential diagnostic information is more important than achieving high n-gram overlap with arbitrary reference texts. These methodological contributions, in combination with the strong empirical results, position this work as a meaningful step toward the eventual integration of generative AI assistants into routine radiological practice.

9.1 Future Scope

Despite the tremendous strides encapsulated in this project, significant avenues for future enhancement remain open:

1. **Large-Scale Benchmark Scaling:** The immediate subsequent phase requires expanding the evaluation set from 50 images to the colossal cohorts found in standard benchmark repositories like MIMIC-CXR and OpenI to truly battle-test the parameters against profound edge cases.
2. **Multi-View Architecture:** Chest radiology routinely requires lateral viewpoints alongside frontal scans to accurately judge cardiac enlargement and posterior lung bases. Upgrading the visual encoder to dynamically parse and fuse multiple input angles simultaneously is a critical next step.
3. **Explainability and Heatmaps:** To foster supreme trust among medical practitioners, it is imperative to integrate visual explainability. Future updates should backward-pass the cross-attention weights to project heatmaps directly over the source image, visually highlighting the precise optical nodes the AI fixated upon when deciding to output words like 'opacity'.
4. **Federated Learning for Privacy-Preserving Training:** One of the most significant barriers to improving medical AI systems is the difficulty of aggregating training data across institutional boundaries due to patient privacy regulations. Implementing a federated learning framework would enable the model to benefit from the diverse pathology distributions present across multiple hospitals without requiring the centralization of sensitive patient data. Each participating institution would train a local

model on its proprietary dataset, and only the aggregated gradient updates would be shared, preserving patient confidentiality while dramatically expanding the effective training corpus.

5. Real-Time Collaborative Annotation: Integrating a feedback loop whereby radiologists can directly edit, annotate, and approve or reject generated reports within the web interface would serve a dual purpose. Clinically, it would transform the system from a standalone generator into a collaborative drafting tool that accelerates rather than replaces the radiologist's workflow. Scientifically, the accumulated corrections would constitute a continuously growing fine-tuning dataset that enables the model to adapt to institution-specific reporting preferences, terminology conventions, and diagnostic priorities through periodic retraining cycles.

6. Regulatory Pathway and Clinical Validation: Transitioning this research prototype into a deployable medical device requires navigating the regulatory frameworks established by bodies such as the FDA and CE marking authorities. This involves conducting prospective clinical trials with statistically powered sample sizes, establishing post-market surveillance protocols for monitoring model drift and performance degradation over time, and developing comprehensive risk management documentation in accordance with ISO 14971 standards for medical device risk analysis. Engaging with regulatory science experts early in the development lifecycle will be essential to ensuring that technical design decisions align with the evidentiary standards required for market authorization.

REFERENCES

- [1] X. Wang, Y. Peng, L. Lu, Z. Lu, M. Bagheri, and R. M. Summers, “ChestX-ray8: Hospital-scale chest X-ray database and benchmarks on weakly-supervised classification and localization of common thorax diseases,” in Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR), 2017, pp. 2097–2106.
- [2] P. Rajpurkar, J. Irvin, K. Zhu et al., “CheXNet: Radiologist-level pneumonia detection on chest X-rays with deep learning,” arXiv preprint arXiv:1711.05225, 2017.
- [3] A. Dosovitskiy, L. Beyer, A. Kolesnikov et al., “An image is worth 16x16 words: Transformers for image recognition at scale,” in Proc. Int. Conf. Learn. Represent. (ICLR), 2021.
- [4] J. Irvin, P. Rajpurkar, M. Ko et al., “CheXpert: A large chest radiograph dataset with uncertainty labels and expert comparison,” in Proc. AAAI Conf. Artif. Intell., vol. 33, 2019, pp. 590–597.
- [5] B. Jing, P. Xie, and E. Xing, “On the automatic generation of medical imaging reports,” in Proc. Annu. Meeting Assoc. Comput. Linguist. (ACL), 2018, pp. 2577–2586.
- [6] Z. Chen, Y. Shen, Y. Song, and X. Wan, “Cross-modal memory networks for radiology report generation,” in Proc. Annu. Meeting Assoc. Comput. Linguist. (ACL), 2022, pp. 5904–5914.
- [7] M. Caron, H. Touvron, I. Misra et al., “Emerging properties in self supervised vision transformers,” in Proc. IEEE Int. Conf. Comput. Vis. (ICCV), 2021, pp. 9650–9660.
- [8] K. Papineni, S. Roukos, T. Ward, and W.-J. Zhu, “BLEU: a method for automatic evaluation of machine translation,” in Proc. Annu. Meeting Assoc. Comput. Linguist. (ACL), 2002, pp. 311–318.
- [9] C.-Y. Lin, “ROUGE: A package for automatic evaluation of summaries,” in Text Summarization Branches Out, 2004, pp. 74–81.

- [10] J. P. Cohen, P. Morrison, L. Dao et al., “COVID-19 image data collection: Prospective predictions are the future,” arXiv preprint arXiv:2006.11988, 2020.
- [11] R. Luo, L. Sun, Y. Xia, T. Qin, S. Zhang, H. Poon, and T.-Y. Liu, “BioGPT: Generative pre-trained transformer for biomedical text generation and mining,” *Briefings in Bioinformatics*, vol. 23, no. 6, 2022.
- [12] G. Litjens, T. Kooi, B. E. Bejnordi et al., “A survey on deep learning in medical image analysis,” *Medical Image Analysis*, vol. 42, pp. 60–88, 2017.
- [13] A. Esteva, A. Robicquet, B. Ramsundar et al., “A guide to deep learning in healthcare,” *Nature Medicine*, vol. 25, no. 1, pp. 24–29, 2019.
- [14] E. J. Topol, “High-performance medicine: the convergence of human and artificial intelligence,” *Nature Medicine*, vol. 25, no. 1, pp. 44–56, 2019.
- [15] A. Hosny, C. Parmar, J. Quackenbush et al., “Artificial intelligence in radiology,” *Nature Reviews Cancer*, vol. 18, no. 8, pp. 500–510, 2018.

Appendix A: Core Implementation Source Code

This appendix contains the core source code for the Medical Report Generator (V5), encompassing the multimodal model architecture (ViT + BioGPT) and the Flask web service.

app.py (Flask Backend & Model Definition Worker)

Medical Report Generator — Flask Web App (V5)

=====

Serves the V5 multimodal model (ViT + BioGPT) for generating chest X-ray radiology reports via a web interface.

```
import io
import os
import re
import time
import traceback
import numpy as np
import torch
import torch.nn as nn

from flask import Flask, jsonify, request, send_from_directory
from PIL import Image
from transformers import AutoModelForCausalLM, AutoTokenizer, AutoModel

# ===== CONFIG =====

CHECKPOINT_PATH = os.path.join(os.path.dirname(__file__),
                                'version_5', 'best_multimodal_v5.pt')

IMAGE_SIZE = 224

IMAGE_MEAN = [0.485, 0.456, 0.406]
```

```
IMAGE_STD = [0.229, 0.224, 0.225]
```

```
# Generation
```

```
NUM_BEAMS = 4
```

```
LENGTH_PENALTY = 1.2
```

```
REPETITION_PENALTY = 2.5
```

```
NO_REPEAT_NGRAM = 3
```

```
MIN_GEN_TOKENS = 50
```

```
MAX_GEN_TOKENS = 256
```

```
# Cross-attention (must match training)
```

```
CROSS_ATTN_HEADS = 8
```

```
CROSS_ATTN_DROPOUT = 0.1
```

```
# ===== MODEL ARCHITECTURE =====
```

```
class CrossAttentionFusion(nn.Module):
```

```
    def __init__(self, hidden_size, num_heads=CROSS_ATTN_HEADS,  
                  dropout=CROSS_ATTN_DROPOUT):
```

```
        super().__init__()
```

```
        self.cross_attn = nn.MultiheadAttention(  
            hidden_size, num_heads, dropout=dropout, batch_first=True)
```

```
        self.norm1 = nn.LayerNorm(hidden_size)
```

```
        self.ffn = nn.Sequential(  
            nn.Linear(hidden_size, hidden_size * 4), nn.GELU(),  
            nn.Dropout(dropout),  
            nn.Linear(hidden_size * 4, hidden_size), nn.Dropout(dropout))
```

```
        self.norm2 = nn.LayerNorm(hidden_size)
```

```
    def forward(self, text_emb, img_emb):
```

```
        attn_out, _ = self.cross_attn(query=text_emb, key=img_emb, value=img_emb)
```

```

        text_emb = self.norm1(text_emb + attn_out)
        return self.norm2(text_emb + self.ffn(text_emb))

class BioGPTMultimodal(nn.Module):
    def __init__(self, vit_model, gpt_model):
        super().__init__()
        self.vit = vit_model
        self.gpt = gpt_model
        gpt_hidden = gpt_model.config.hidden_size
        self.proj = nn.Linear(vit_model.config.hidden_size, gpt_hidden)
        self.cross_attn = CrossAttentionFusion(gpt_hidden)

    def forward(self, pixel_values, input_ids, attention_mask, labels=None):
        vit_out = self.vit(pixel_values=pixel_values).last_hidden_state
        img_emb = self.proj(vit_out)
        token_emb = self.gpt.get_input_embeddings()(input_ids)
        fused_text = self.cross_attn(token_emb, img_emb)
        inputs_embeds = torch.cat([img_emb, fused_text], dim=1)
        img_mask = torch.ones((input_ids.size(0), img_emb.size(1)),
                               dtype=attention_mask.dtype,
                               device=attention_mask.device)
        new_attn = torch.cat([img_mask, attention_mask], dim=1)
        new_labels = None
        if labels is not None:
            pad = torch.full((labels.size(0), img_emb.size(1)), -100,
                              dtype=labels.dtype, device=labels.device)
            new_labels = torch.cat([pad, labels], dim=1)
        return self.gpt(inputs_embeds=inputs_embeds,
                        attention_mask=new_attn, labels=new_labels,

```

```
        return_dict=True)

# ===== INFERENCE =====

def preprocess_image(pil_image, image_size=IMAGE_SIZE):
    """Preprocess a PIL image for the ViT model."""
    mean = np.array(IMAGE_MEAN, dtype='float32').reshape(1, 1, 3)
    std = np.array(IMAGE_STD, dtype='float32').reshape(1, 1, 3)
    im = pil_image.convert('RGB').resize((image_size, image_size),
                                         resample=Image.BILINEAR)
    arr = np.array(im).astype('float32') / 255.0
    arr = (arr - mean) / std
    arr = np.transpose(arr, (2, 0, 1))
    return torch.from_numpy(arr).unsqueeze(0).to(torch.float32)

def generate_report(model, tokenizer, pil_image, history_text="",
                   device='cuda'):
    """Generate a raw report from a PIL image."""
    model.eval()
    pix = preprocess_image(pil_image).to(device)
    prompt = f'Patient history: {history_text}\nReport: '
    prefix_enc = tokenizer(prompt, return_tensors='pt').to(device)
    prefix_ids = prefix_enc['input_ids']
    prefix_attn = prefix_enc['attention_mask']
    with torch.no_grad():
        vit_out = model.vit(pixel_values=pix).last_hidden_state
        img_emb = model.proj(vit_out)
        token_emb = model.gpt.get_input_embeddings()(prefix_ids)
```

```

fused_text = model.cross_attn(token_emb, img_emb)

inputs_embeds = torch.cat([img_emb, fused_text], dim=1)
img_mask = torch.ones((1, img_emb.size(1)),
                       dtype=prefix_attn.dtype).to(device)
new_mask = torch.cat([img_mask, prefix_attn], dim=1)
gen_ids = model.gpt.generate(
    inputs_embeds=inputs_embeds,
    attention_mask=new_mask,
    max_new_tokens=MAX_GEN_TOKENS,
    min_new_tokens=MIN_GEN_TOKENS,
    num_beams=NUM_BEAMS,
    length_penalty=LENGTH_PENALTY,
    repetition_penalty=REPETITION_PENALTY,
    no_repeat_ngram_size=NO_REPEAT_NGRAM,
)
out_text = tokenizer.decode(gen_ids[0], skip_special_tokens=True)
if prompt in out_text:
    out_text = out_text.split(prompt, 1)[1]
return out_text.strip()

def smart_sentence_split(text):
    abbreviations = (r'(<!Dr)(<!Mr)(<!Mrs)(<!Ms)(<!No)'
                     r'(<!vs)(<!etc)(<!i\\.e)(<!e\\.g)')
    sentences = re.split(abbreviations + r'\\.s+', text)
    return [s.strip() for s in sentences if s.strip()]

def parse_report_sections(raw_text):
    # ... omitted string formatting logic for brevity ...

```

```
    return {"findings": [], "impression": [], "recommendations": []}

# ===== FLASK APP =====

from flask_cors import CORS
app = Flask(__name__, static_folder='static')
CORS(app)
# Globals for model
model = None
tokenizer = None
device = None
def load_model():
    """Load the V5 model at startup."""
    global model, tokenizer, device
    if torch.cuda.is_available():
        device = 'cuda'
        props = torch.cuda.get_device_properties(0)
        print(f' GPU: {props.name} ( {props.total_memory / 1e9:.1f} GB VRAM)')
    else:
        device = 'cpu'
        print(' No GPU detected — using CPU (slower)')
    print('\n[1/4] Loading tokenizer...')
    tokenizer = AutoTokenizer.from_pretrained('microsoft/biogpt')
    if tokenizer.pad_token is None:
        tokenizer.add_special_tokens({'pad_token': '<pad>'})
    print('[2/4] Loading ViT (microsoft/rad-dino)...')
    try:
```



```
vit = AutoModel.from_pretrained('microsoft/rad-dino')
print(' ✓ rad-dino loaded')
except Exception:
    print(' rad-dino failed, falling back to vit-base-patch16-224-in21k')
    from transformers import ViTModel
    vit = ViTModel.from_pretrained('google/vit-base-patch16-224-in21k')
print('[3/4] Loading BioGPT...')
biogpt = AutoModelForCausalLM.from_pretrained('microsoft/biogpt')
biogpt.resize_token_embeddings(len(tokenizer))
print('[4/4] Loading V5 checkpoint...')
model = BioGPTMultimodal(vit, biogpt)
ckpt = torch.load(CHECKPOINT_PATH, map_location=device, weights_only=False)
result = model.load_state_dict(ckpt['model_state_dict'], strict=False)
if result.missing_keys:
    print(f' Note: {len(result.missing_keys)} missing keys (expected for architecture diff)')
    if result.unexpected_keys:
        print(f' Note: {len(result.unexpected_keys)} unexpected keys')
print(' ✓ Checkpoint loaded successfully')
model = model.to(device)
print(f' Model loaded in FP32 on {device}')
model.eval()
total = sum(p.numel() for p in model.parameters())
print(f'\n 🟢 Model ready! ({total:,} parameters)')
print(f' Device: {device}')
print(f' Open http://localhost:5000 in your browser\n')
@app.route('/')
```

```
def index():
    return send_from_directory('static', 'index.html')

@app.route('/api/generate', methods=['POST'])
def api_generate():
    """Generate a report from an uploaded X-ray image."""
    if model is None:
        return jsonify({'error': 'Model not loaded yet'}), 503
    if 'image' not in request.files:
        return jsonify({'error': 'No image uploaded'}), 400
    file = request.files['image']
    if file.filename == "":
        return jsonify({'error': 'No file selected'}), 400
    history = request.form.get('history', "")
    try:
        pil_image = Image.open(io.BytesIO(file.read()))
        start = time.time()
        raw = generate_report(model, tokenizer, pil_image,
                              history_text=history, device=device)
        elapsed = time.time() - start
        sections = parse_report_sections(raw)
        return jsonify({
            'raw_report': raw,
            'sections': sections,
            'inference_time': round(elapsed, 2),
            'device': device,
        })
    except Exception as e:
```

```
        traceback.print_exc()

        return jsonify({'error': str(e)}), 500

@app.route('/api/health')
def health():

    return jsonify({

        'status': 'ready' if model is not None else 'loading',

        'device': device or 'unknown',

    })

if __name__ == '__main__':

    print('=' * 60)

    print(' Medical Report Generator — V5')

    print(' ViT (rad-dino) + BioGPT + Cross-Attention')

    print('=' * 60)

    load_model()

    app.run(host='0.0.0.0', port=5000, debug=False)
```